

```
In [308]: import os
import pandas as pd
import matplotlib
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import pickle
from sklearn.manifold import TSNE
from sklearn import preprocessing
import pandas as pd
```

Initial data processing

```
In [309]: data=pd.read_csv("C://Users//sthakal//OneDrive - George Weston Limited-6469347-MTCAD//Psnal/
```

```
In [310]: data.head(15)
```

Out[310]:

	age	job	marital	education	default	housing	loan	contact	month	day_of_week	...	ca
0	56	housemaid	married	basic.4y	no	no	no	telephone	may	mon	...	
1	57	services	married	high.school	unknown	no	no	telephone	may	mon	...	
2	37	services	married	high.school	no	yes	no	telephone	may	mon	...	
3	40	admin.	married	basic.6y	no	no	no	telephone	may	mon	...	
4	56	services	married	high.school	no	no	yes	telephone	may	mon	...	
5	45	services	married	basic.9y	unknown	no	no	telephone	may	mon	...	
6	59	admin.	married	professional.course	no	no	no	telephone	may	mon	...	
7	41	blue-collar	married	unknown	unknown	no	no	telephone	may	mon	...	
8	24	technician	single	professional.course	no	yes	no	telephone	may	mon	...	
9	25	services	single	high.school	no	yes	no	telephone	may	mon	...	
10	41	blue-collar	married	unknown	unknown	no	no	telephone	may	mon	...	
11	25	services	single	high.school	no	yes	no	telephone	may	mon	...	
12	29	blue-collar	single	high.school	no	no	yes	telephone	may	mon	...	
13	57	housemaid	divorced	basic.4y	no	yes	no	telephone	may	mon	...	
14	35	blue-collar	married	basic.6y	no	yes	no	telephone	may	mon	...	

15 rows × 21 columns

```
In [311]: data.dtypes
```

```
Out[311]: age                int64
job                object
marital           object
education         object
default          object
housing          object
loan             object
contact          object
month            object
day_of_week      object
duration         int64
campaign         int64
pdays          int64
previous         int64
poutcome         object
emp.var.rate     float64
cons.price.idx   float64
cons.conf.idx    float64
euribor3m        float64
nr.employed      float64
y               object
dtype: object
```

```
In [312]: data.isna().count()
# CHECKING missing values
```

```
Out[312]: age                41188
job                41188
marital           41188
education         41188
default          41188
housing          41188
loan             41188
contact          41188
month            41188
day_of_week      41188
duration         41188
campaign         41188
pdays          41188
previous         41188
poutcome         41188
emp.var.rate     41188
cons.price.idx   41188
cons.conf.idx    41188
euribor3m        41188
nr.employed      41188
y               41188
dtype: int64
```

```
In [313]: data.shape
```

```
Out[313]: (41188, 21)
```


```
In [314]: data.describe()
#Checking the stats of numerical variable
```

Out[314]:

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.c
count	41188.00000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.000000	41188.
mean	40.02406	258.285010	2.567593	962.475454	0.172963	0.081886	93.575664	-40.
std	10.42125	259.279249	2.770014	186.910907	0.494901	1.570960	0.578840	4.
min	17.00000	0.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.
25%	32.00000	102.000000	1.000000	999.000000	0.000000	-1.800000	93.075000	-42.
50%	38.00000	180.000000	2.000000	999.000000	0.000000	1.100000	93.749000	-41.
75%	47.00000	319.000000	3.000000	999.000000	0.000000	1.400000	93.994000	-36.
max	98.00000	4918.000000	56.000000	999.000000	7.000000	1.400000	94.767000	-26.

```
In [315]: data.hist('age')
```

Out[315]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000026732009E8>]],
dtype=object)



Checking the counts of Categorical variables

In [316]: *#we are now seeing the counts of the categorical variable one by one to see the values*
#Checking the counts of categorical variables to see what value they are concentrated in

```
job_cnt=data['job'].value_counts()  
print(job_cnt)
```

```
admin.          10422  
blue-collar     9254  
technician      6743  
services        3969  
management     2924  
retired         1720  
entrepreneur    1456  
self-employed   1421  
housemaid       1060  
unemployed      1014  
student         875  
unknown         330  
Name: job, dtype: int64
```

In [317]: data['marital'].value_counts()

```
Out[317]: married      24928  
single      11568  
divorced     4612  
unknown        80  
Name: marital, dtype: int64
```

In [318]: data['education'].value_counts()

```
Out[318]: university.degree    12168  
high.school      9515  
basic.9y         6045  
professional.course  5243  
basic.4y         4176  
basic.6y         2292  
unknown         1731  
illiterate        18  
Name: education, dtype: int64
```

In [319]: data['marital'].value_counts()

```
Out[319]: married      24928  
single      11568  
divorced     4612  
unknown        80  
Name: marital, dtype: int64
```

In [320]: data['education'].value_counts()

```
Out[320]: university.degree    12168  
high.school      9515  
basic.9y         6045  
professional.course  5243  
basic.4y         4176  
basic.6y         2292  
unknown         1731  
illiterate        18  
Name: education, dtype: int64
```

```
In [321]: data['housing'].value_counts()
```

```
Out[321]: yes          21576  
         no           18622  
         unknown       990  
         Name: housing, dtype: int64
```

```
In [322]: data['loan'].value_counts()
```

```
Out[322]: no          33950  
         yes          6248  
         unknown       990  
         Name: loan, dtype: int64
```

```
In [323]: data['contact'].value_counts()
```

```
Out[323]: cellular     26144  
         telephone    15044  
         Name: contact, dtype: int64
```

```
In [324]: data['month'].value_counts()
```

```
Out[324]: may          13769  
         jul           7174  
         aug           6178  
         jun           5318  
         nov           4101  
         apr           2632  
         oct            718  
         sep            570  
         mar            546  
         dec            182  
         Name: month, dtype: int64
```

```
In [325]: data['day_of_week'].value_counts()
```

```
Out[325]: thu          8623  
         mon          8514  
         wed          8134  
         tue          8090  
         fri          7827  
         Name: day_of_week, dtype: int64
```

```
In [326]: data['poutcome'].value_counts()
```

```
Out[326]: nonexistent   35563  
         failure         4252  
         success         1373  
         Name: poutcome, dtype: int64
```

```
In [327]: data['month'].value_counts()
```

```
Out[327]: may      13769  
jul       7174  
aug       6178  
jun       5318  
nov       4101  
apr       2632  
oct        718  
sep        570  
mar        546  
dec        182  
Name: month, dtype: int64
```

```
In [328]: data['y'].value_counts()  
#understanding the count of the Predictor variable
```

```
Out[328]: no      36548  
yes      4640  
Name: y, dtype: int64
```


```
In [329]: data['default'].value_counts()
```

```
Out[329]: no      32588  
unknown   8597  
yes        3  
Name: default, dtype: int64
```

Bivariate Analysis

```
In [330]: data.plot.scatter('euribor3m', 'emp.var.rate',)
```

```
Out[330]: <matplotlib.axes._subplots.AxesSubplot at 0x26733590048>
```



```
In [331]: data.corr()
#Checking numerical correlation
```

```
Out[331]:
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.idx
age	1.000000	-0.000866	0.004594	-0.034369	0.024365	-0.000371	0.000857	0.129372
duration	-0.000866	1.000000	-0.071699	-0.047577	0.020640	-0.027968	0.005312	-0.008173
campaign	0.004594	-0.071699	1.000000	0.052584	-0.079141	0.150754	0.127836	-0.013733
pdays	-0.034369	-0.047577	0.052584	1.000000	-0.587514	0.271004	0.078889	-0.091342
previous	0.024365	0.020640	-0.079141	-0.587514	1.000000	-0.420489	-0.203130	-0.050936
emp.var.rate	-0.000371	-0.027968	0.150754	0.271004	-0.420489	1.000000	0.775334	0.196041
cons.price.idx	0.000857	0.005312	0.127836	0.078889	-0.203130	0.775334	1.000000	0.058986
cons.conf.idx	0.129372	-0.008173	-0.013733	-0.091342	-0.050936	0.196041	0.058986	1.000000
euribor3m	0.010767	-0.032897	0.135133	0.296899	-0.454494	0.972245	0.688230	0.277686
nr.employed	-0.017725	-0.044703	0.144095	0.372605	-0.501333	0.906970	0.522034	0.100513

Some preprocessing for Bivariate

```
In [332]: #Clubbing all basic education in one
data['education']=np.where(data['education']=='basic.9y', 'Basic', data['education'])
data['education']=np.where(data['education']=='basic.6y', 'Basic', data['education'])
data['education']=np.where(data['education']=='basic.4y', 'Basic', data['education'])
```

```
In [333]: data['y']=np.where(data['y']=='yes',1, data['y'])# changing into binary/
data['y']=np.where(data['y']=='no',0, data['y'])
```


```
In [334]: #Checking how the categorical variable is related to the Response variable
%matplotlib inline

def countplot_withY(label, dataset):
    plt.figure(figsize=(20,10))
    Y = data[label]
    total = len(Y)*1.
    ax=sns.countplot(x=label, data=dataset, hue="y")
    for p in ax.patches:
        ax.annotate('{:.1f}%'.format(100*p.get_height()/total), (p.get_x()+0.1, p.get_height()+5


    #put 11 ticks (therefore 10 steps), from 0 to the total number of rows in the dataframe
    ax.yaxis.set_ticks(np.linspace(0, total, 11))
    #adjust the ticklabel to the desired format, without changing the position of the ticks.
    ax.set_yticklabels(map('{:.1f}%'.format, 100*ax.yaxis.get_majorticklocs()/total))
    ax.set_xticklabels(ax.get_xticklabels(), rotation=40, ha="right")
    # ax.legend(labels=["no", "yes"])
    plt.show()
```

Seeing the relationship of categorical variables with Response Variable


```
In [335]: countplot_withY("marital",data)
```




```
In [336]: countplot_withY("education",data)
```





```
In [337]: countplot_withY("housing",data)
```




```
In [338]: countplot_withY("loan",data)
```




```
In [339]: countplot_withY("contact",data)
```




```
In [340]: countplot_withY("month",data)
```




```
In [341]: countplot_withY("day_of_week",data)
```



```
In [342]: countplot_withY("nr.employed",data)
```



```
In [343]: countplot_withY("poutcome", data)
```



```
In [344]: data_response_yes=data[data['y']==1]
#only selecting those dataset where response is yes
```

```
In [345]: data_response_yes.describe()
#Just checking how the numerical variable of success looks like
```

```
Out[345]:
```

	age	duration	campaign	pdays	previous	emp.var.rate	cons.price.idx	cons.conf.
count	4640.000000	4640.000000	4640.000000	4640.000000	4640.000000	4640.000000	4640.000000	4640.000000
mean	40.913147	553.191164	2.051724	792.035560	0.492672	-1.233448	93.354386	-39.7891
std	13.837476	401.171871	1.666245	403.407181	0.860344	1.623626	0.676644	6.1396
min	17.000000	37.000000	1.000000	0.000000	0.000000	-3.400000	92.201000	-50.8000
25%	31.000000	253.000000	1.000000	999.000000	0.000000	-1.800000	92.893000	-46.2000
50%	37.000000	449.000000	2.000000	999.000000	0.000000	-1.800000	93.200000	-40.4000
75%	50.000000	741.250000	2.000000	999.000000	1.000000	-0.100000	93.918000	-36.1000
max	98.000000	4199.000000	23.000000	999.000000	6.000000	1.400000	94.767000	-26.9000

Some preprocessing for model building

```
In [346]: #Create dummy variables
data_x = data.iloc[:, :-1]
print("Shape of X:", data_x.shape)
data_y = data["y"]
print("Shape of Y:", data_y.shape)
```

Shape of X: (41188, 20)

Shape of Y: (41188,)

```
In [347]: from sklearn.model_selection import train_test_split

X_rest, X_test, y_rest, y_test = train_test_split(data_x, data_y, test_size=0.2)
X_train, X_cv, y_train, y_cv = train_test_split(X_rest, y_rest, test_size=0.2)

print("X Train:", X_train.shape)
print("X CV:", X_cv.shape)
print("X Test:", X_test.shape)
print("Y Train:", y_train.shape)
print("Y CV:", y_cv.shape)
print("Y Test:", y_test.shape)
```

X Train: (26360, 20)

X CV: (6590, 20)

X Test: (8238, 20)

Y Train: (26360,)

Y CV: (6590,)

Y Test: (8238,)

```
In [348]: y_train.replace({"no":0, "yes":1}, inplace=True)
y_cv.replace({"no":0, "yes":1}, inplace=True)
y_test.replace({"no":0, "yes":1}, inplace=True)
```

```
In [349]: # Categorical boolean mask
categorical_feature_mask = data_x.dtypes==object

# filter categorical columns using mask and turn it into a list
categorical_cols = data_x.columns[categorical_feature_mask].tolist()
```

```
In [350]: data_x.dtypes==object  
#checking object type
```

```
Out[350]: age                False  
job                True  
marital            True  
education          True  
default            True  
housing            True  
loan               True  
contact            True  
month              True  
day_of_week        True  
duration            False  
campaign            False  
pdays             False  
previous            False  
poutcome            True  
emp.var.rate       False  
cons.price.idx     False  
cons.conf.idx      False  
euribor3m          False  
nr.employed        False  
dtype: bool
```

```

In [351]: from sklearn.feature_extraction.text import CountVectorizer

def add_onehot_to_dataframe(sparse, df, vectorizer, name):
    """
        This function will add the one hot encoded to the dataframe.
    """
    for i, col in enumerate(vectorizer.get_feature_names()):
        colname = name+"_"+col
        # df[colname] = pd.SparseSeries(sparse[:, i].toarray().flatten(), fill_value=0)
        df[colname] = sparse[:, i].toarray().ravel().tolist()

    return df

def OneHotEncoder(categorical_cols, X_train, X_test, X_cv=None, include_cv=False):
    """
        This function takes categorical column names as inputs. The objective
        of this function is to take the column names iteratively and encode the
        features using One hot Encoding mechanism and also adding the encoded feature
        to the respective dataframe.

        The include_cv parameter indicates whether we should include CV dataset or not.
        This is added specifically because when using GridSearchCV or RandomizedSearchCV,
        we only split the dataset into train and test to give more data to training purposes.
        This is done because GridSearchCV splits the data internally anyway.
    """

    for i in categorical_cols:
        Vectorizer = CountVectorizer(token_pattern="[A-Za-z0-9-.-]+")
        print("Encoding for feature: ", i)
        # Encoding training dataset
        temp_cols = Vectorizer.fit_transform(X_train[i])
        X_train = add_onehot_to_dataframe(temp_cols, X_train, Vectorizer, i)

        # Encoding Cross validation dataset
        if include_cv:
            temp_cols = Vectorizer.transform(X_cv[i])
            X_cv = add_onehot_to_dataframe(temp_cols, X_cv, Vectorizer, i)

        # Encoding Test dataset
        temp_cols = Vectorizer.transform(X_test[i])
        X_test = add_onehot_to_dataframe(temp_cols, X_test, Vectorizer, i)

```

In [352]: `OneHotEncoder(categorical_cols, X_train, X_test, X_cv, True)`

```
# Drop the categorical features as the one hot encoded representation is present
X_train = X_train.drop(categorical_cols, axis=1)
X_cv = X_cv.drop(categorical_cols, axis=1)
X_test = X_test.drop(categorical_cols, axis=1)

print("Shape of train: ", X_train.shape)
print("Shape of CV: ", X_cv.shape)
print("Shape of test: ", X_test.shape)
```

Encoding for feature: job

C:\Users\sthakal\AppData\Local\Continuum\anaconda3\lib\site-packages\ipykernel_launcher.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

`# This is added back by InteractiveShellApp.init_path()`

Encoding for feature: marital
Encoding for feature: education
Encoding for feature: default
Encoding for feature: housing
Encoding for feature: loan
Encoding for feature: contact
Encoding for feature: month
Encoding for feature: day of week

In [356]: `# with "duration" column`
`from sklearn.metrics import roc_auc_score`
`from sklearn.linear_model import LogisticRegression`

```
model = LogisticRegression(class_weight='balanced',max_iter=500)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)

print("AUC score with duration column: ", roc_auc_score(y_test, y_pred[:,1]))
```

AUC score with duration column: 0.7902335490207407

In [354]: `# Removing duration feature`

```
# From Train
X_train = X_train.drop("duration", axis=1)
print("The shape of the train dataset: ", X_train.shape)

# From CV
X_cv = X_cv.drop("duration", axis=1)
print("The shape of the cv dataset: ", X_cv.shape)

# From Test
X_test = X_test.drop("duration", axis=1)
print("The shape of the test dataset: ", X_test.shape)
```

The shape of the train dataset: (26360, 60)
The shape of the cv dataset: (6590, 60)
The shape of the test dataset: (8238, 60)


```
In [357]: # without "duration" column

model = LogisticRegression(class_weight='balanced',max_iter=500)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)

print("AUC score without duration column: ", roc_auc_score(y_test, y_pred[:,1]))
```

AUC score without duration column: 0.7902335490207407

```
In [358]: # without "duration" column
# X_train = X_train.drop("duration", axis=1)
# X_test = X_test.drop("duration", axis=1)

# print(X_train.shape)
# print(X_test.shape)

model = LogisticRegression(class_weight='balanced',max_iter=500)
model.fit(X_train, y_train)
y_pred = model.predict_proba(X_test)

print("AUC score without duration column: ", roc_auc_score(y_test, y_pred[:,1]))
```

AUC score without duration column: 0.7889597897047779